

Thread Safety



Agenda

- Highlight issues with multi threaded programming
- Introduce thread synchronization primitives
- Introduce thread safe collections



Need for Synchronization

- ⬡ Creating threads is easy
- ⬡ When threads share data problems can occur
 - ⬢ Inconsistent reads
 - ⬢ State corruption
- ⬡ Synchronization fixes these problems, but potentially creates a new problem
 - ⬢ Over synchronization reduces scalability
- ⬡ Lots of techniques to implement synchronization
 - ⬢ Each have cost and benefit
- ⬡ Developers role is to write an application that scales and is thread safe, by selecting the best synchronization technique



Simple Increment

Two threads

- Sharing an instance of **Counter**.
- Both are calling **Increment** 1000 times

Question

- What is the value of count after both threads have completed?

```
public class Counter
{
    protected int count;

    public virtual void Increment()
    {
        count++;
    }
    public int Value { get { return count; } }
}
```



Simple Increment, NOT Atomic

- Even a simple `count++` is not an atomic operation.
 - Multiple CPU instructions that could be interweaved.
- Consider the possible execution below of two threads (T0, T1)
 - Assuming `count=0` at the start
 - At the end of execution `i` would be 1 and not the desired 2.
- If two threads don't attempt to increment count at the same time not a problem. Spotting these kind of errors is hard

time
↓

T0: MOV R0, `count`

T0: ADD R0, 1

T0: MOV `count`, R0

T1: MOV R0, `count`

T1: ADD R0, 1

T1: MOV `count`, R0

Interlocked

- Modern CPU's expose special instruction set to perform various operations atomically
 - Cost more than non atomic variants.
- Access to these instructions via **Interlocked** class
 - Interlocked.Increment
 - Interlocked.Decrement
 - Interlocked.Add

```
public class InterlockedCounter : Counter
{
    public override void Increment()
    {
        // Atomic count++
        Interlocked.Increment(ref count);
    }
}
```

Multi step state transition

What happens if

- Thread A is inside ReceivePayment
- Thread B is **inside NetWorth**

Can Interlocked help ?

```
class SmallBusiness {  
    decimal Cash = 0;  
    decimal Receivables = 1000;  
  
    public void ReceivePayment(decimal amount) {  
        Cash += amount;  
        Receivables -= amount;  
    }  
  
    public decimal NetWorth {  
        get { return Cash + Receivables; }  
    }  
}
```



Sequential access

🟡 To fix the problem

- 🟡 Sequentialise access to the object state

🟡 How

- 🟡 Each instance of a reference type has a Monitor
 - 🟡 CLR guarantees that only one thread can own the monitor
 - 🟡 If a thread can't acquire the monitor it enters a wait state
 - 🟡 When the monitor is available it is woken up and proceeds
- 🟡 Critical areas of code can therefore be protected by using a monitor.



Monitor based solution

- Only one thread in any **critical region** at any point in time

```
private object _lock = new object();

public void ReceivePayment(decimal amount)
{
    Monitor.Enter(_lock);
    Cash += amount;
    Receivables -= amount;
    Monitor.Exit(_lock);
}
```

Could be an issue with exceptions

```
public decimal NetWorth
{
    get {
        Monitor.Enter(_lock);
        try { return Cash + Receivables; }
        finally { Monitor.Exit(_lock); }
    }
}
```

Deals better with exceptions



Lock keyword

- Enter, try, finally , Exit common pattern
 - C# language offers **lock** keyword to assist
 - Compiler emits try, finally logic
- Use of Monitor.Enter and **lock** can lead to deadlocks
 - Prefer Montior.TryEnter which takes a timeout
- Avoid using lock(this) and lock(typeof(X))
 - Less control over objects use for synchronization.

```
public void ReceivePayment(decimal amount) {  
    lock (_lock)  
    {  
        Cash += amount;  
        Receivables -= amount;  
    }  
}
```

High Read to Write Ratio

- Monitor provides mutual exclusion behaviour
 - Excluding readers and writers
- Thread safety not an issue if all threads read.
- Better throughput may be achieved with a Synchronization primitive that ensures
 - There can be Many Readers, Zero Writer
 - Or One Writer, Zero Readers
- This is known as a ReaderWriterLock
 - .NET 3.5 and above prefer ReaderWriterLockSlim
 - Pre 3.5, ReaderWriterLock
 - Not well implemented, can result in writer being denied access for long periods of time.
- Often used for caching, where most of the time is spent reading with occasional updates.



Reader Writer Lock

```
public class SimpleCache
{
    private Dictionary<int, string> cache=
        new Dictionary<int, string>();
    private ReaderWriterLockSlim _lock = new ReaderWriterLockSlim();
```

```
    public string Get(int key) {
        _lock.EnterReadLock();
        try { return cache[key]; }
        finally { _lock.ExitReadLock(); }
    }
```

```
    public void Set(int key, string val)
    {
        _lock.EnterWriteLock();
        try { cache.Add(key, val) }
        finally { _lock.ExitWriteLock(); }
    }
}
```

Many threads can
can read from the
cache

When one thread
has the write lock
no other thread
can obtain read or
write lock.

Synchronization across app domains

- **Managed synchronization primitives only allow synchronization inside a single app domain**
- **How to control access to a shared file ?**
 - Requires Kernel based synchronization
- **Kernel synchronization can be achieved via managed wrappers**
 - Mutex
 - Semaphore
 - AutoResetEvent
 - ManualResetEvent
- **These synchronization primitives are orders of magnitude more expensive than managed ones**



Summary

- ⬢ A variety of ways to perform synchronization, the skill is picking the correct one
- ⬢ Concurrent collections make it simpler to write efficient thread safe code
- ⬢ Only use kernel synchronization primitives when absolutely necessary
- ⬢ Analyse code and imagine worse possible race conditions

